

```
# Importing libraries
!pip install -q scikit-learn transformers torch

import pandas as pd
import numpy as np
import torch
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score
from transformers import pipeline
```

```
↳ _____ 363.4/363.4 MB 3.6 MB/s eta 0:00:00
_____ 13.8/13.8 MB 38.5 MB/s eta 0:00:00
_____ 24.6/24.6 MB 29.5 MB/s eta 0:00:00
_____ 883.7/883.7 kB 29.4 MB/s eta 0:00:00
_____ 664.8/664.8 MB 3.0 MB/s eta 0:00:00
_____ 211.5/211.5 MB 5.7 MB/s eta 0:00:00
_____ 56.3/56.3 MB 13.9 MB/s eta 0:00:00
_____ 127.9/127.9 MB 7.6 MB/s eta 0:00:00
_____ 207.5/207.5 MB 6.2 MB/s eta 0:00:00
_____ 21.1/21.1 MB 106.5 MB/s eta 0:00:00
```

```
# Load csv file
df = pd.read_csv('/content/IMDB Dataset.csv')
```

```
df.head()
```

```
↳
```

	review	sentiment
0	One of the other reviewers has mentioned that ...	positive
1	A wonderful little production. The...	positive
2	I thought this was a wonderful way to spend ti...	positive
3	Basically there's a family where a little boy ...	negative
4	Petter Mattei's "Love in the Time of Money" is...	positive

```
# Map sentiment to binary labels
df['label'] = df['sentiment'].map({'positive': 1, 'negative': 0})

# Defining features and labels
X = df['review'].values
y = df['label'].values

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

```
# TF-IDF Vectorization
tfidf = TfidfVectorizer(max_features=5000)
X_train_tfidf = tfidf.fit_transform(X_train)
X_test_tfidf = tfidf.transform(X_test)
```

```
# Training Logistic Regression
lr_model = LogisticRegression(max_iter=1000)
lr_model.fit(X_train_tfidf, y_train)
```

```
# Predict
y_pred_lr = lr_model.predict(X_test_tfidf)
```

```
# Evaluation
print("=== Logistic Regression Performance ===")
print(classification_report(y_test, y_pred_lr))
print("Accuracy:", accuracy_score(y_test, y_pred_lr))
```

```
↳ === Logistic Regression Performance ===
```

	precision	recall	f1-score	support
0	0.90	0.89	0.89	5000
1	0.89	0.90	0.89	5000

```

    accuracy      0.89      0.89      0.89      10000
    macro avg      0.89      0.89      0.89      10000
    weighted avg   0.89      0.89      0.89      10000

```

Accuracy: 0.8924

```

# Loading BERT pipeline
bert_classifier = pipeline(
    "sentiment-analysis",
    model="textattack/bert-base-uncased-imdb",
    truncation=True,
    device=0
)

```

```

# Predicting
bert_preds = []
for text in X_test:
    if len(text) > 1000:
        text = text[:1000]

    output = bert_classifier(text)[0]['label']
    label = 1 if output == 'LABEL_1' else 0
    bert_preds.append(label)

```

```

# Evaluating
print("\n=== API-based BERT Performance ===")
print(classification_report(y_test, bert_preds))
print("Accuracy:", accuracy_score(y_test, bert_preds))

```

⚠ /usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>), set
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets.

```

warnings.warn(
    config.json: 100%          511/511 [00:00<00:00, 10.3kB/s]
    pytorch_model.bin: 100%    438M/438M [00:02<00:00, 224MB/s]
    tokenizer_config.json: 100% 48.0/48.0 [00:00<00:00, 5.35kB/s]
    model.safetensors: 100%    438M/438M [00:02<00:00, 152MB/s]
    vocab.txt: 100%            232k/232k [00:00<00:00, 4.08MB/s]
    special_tokens_map.json: 100% 112/112 [00:00<00:00, 7.59kB/s]

```

Device set to use cuda:0
You seem to be using the pipelines sequentially on GPU. In order to maximize efficiency please use a dataset

```

=== BERT Performance ===
      precision    recall  f1-score   support

     0       0.95      0.96      0.95      5000
     1       0.96      0.95      0.95      5000

   accuracy      0.95      0.95      0.95      10000
  macro avg      0.95      0.95      0.95      10000
 weighted avg      0.95      0.95      0.95      10000

```

Accuracy: 0.9523

```

# Comparison summary
results = {
    'Model': ['Logistic Regression', 'API-based BERT'],
    'Accuracy': [
        accuracy_score(y_test, y_pred_lr),
        accuracy_score(y_test, bert_preds)
    ]
}

```

```

results_df = pd.DataFrame(results)
print(results_df)

```

⚠

	Model	Accuracy
0	Logistic Regression	0.8924
1	API-based BERT	0.9523

